

FTML practical session 5

17 mai 2026




TABLE DES MATIÈRES

1	Curse of dimensionality and adaptivity to a low dimensional support	2
1.1	Nearest neighbors algorithms	2
1.2	1D and 2D case	3

1.3	Curse of dimensionality : general case	3
1.4	Adaptivity	4

1 CURSE OF DIMENSIONALITY AND ADAPTIVITY TO A LOW DIMENSIONAL SUPPORT

1.1 Nearest neighbors algorithms

A nearest neighbors algorithm is a method of prediction, based on a dataset, but without any optimization. It consists in averaging the predictions of the nearest neighbors of a sample x in a given dataset. See two examples below, where both input and output spaces are $\mathbb{R}$. The two examples differ by the number of samples in the dataset.




FIGURE 1 – knn




FIGURE 2 – knn

A natural question would be : why do we not use this method all the time? The answer is that as soon as the input space dimension is not very small (for instance not < 10), the number of samples that would be required to obtain a satisfactory error is simply too large.

1.2 1D and 2D case

1.2.1 1D case

We consider a L -Lipshitz-continuous function f , defined over $[0, 1]$.

$$\forall x, y \in [0, 1], |f(x) - f(y)| \leq L|x - y| \quad (1)$$

Show that if we have n samples, we can approximate f with a nearest neighbor method with an error smaller than $\frac{L}{2n}$.

1.2.2 2D case

We consider a L -Lipshitz-continuous function f , now defined over $[0, 1]^2$.

$$\forall x, y \in [0, 1]^2, |f(x) - f(y)| \leq L|x - y| \quad (2)$$

Show that if we have n samples, we can approximate f with a nearest neighbor method with an error smaller than $\frac{L}{\sqrt{2n}}$.

1.3 Curse of dimensionality : general case

In the general case, it is possible to show that if the target function f^* is Lipshitz-continuous, and if the n samples in the dataset are on a lattice that divides $[0, 1]^d$ in cubes of equal size, then the error made by a k -NN estimator can be upper bounded by a bound of the form $\mathcal{O}(n^{-1/d})$.

Note that this is a rough bound. It is possible to have more subtle results, and the constant in the $\mathcal{O}$ depends on d and on the Lipshitz constant of the target function. However, it is not possible to have a bound that is "way better" (i.e. "way smaller") than this bound. This means that if d is large, then $1/d$ is small, and the decrease of the error is in general very slow (meaning : in general, we should expect to have a decrease that is as slow as this one).

With additional hypotheses, it is also possible to show that in order to guarantee an error as small as some value $\epsilon > 0$, the number of samples needed n_ϵ verifies

$$n_\epsilon \geq \frac{\epsilon^{-d} d^{d/2}}{\alpha^{d/2}} \quad (3)$$

where $\alpha > 0$ is a constant. n_ϵ hence grows exponentially with $1/\epsilon$, and more than exponentially with d which is an example of curse of dimensionality.

We note that our target function is indeed Lipshitz continuous and thus respects the hypothesis to obtain equation 3. If a function is not Lipshitz continuous, the situation is even worse!

1.3.1 Simulation : uniformly distributed data in the input space

We will study the influence of the input space dimension on the possibility to approximate a function with a nearest neighbors approach in a specific setting :

- input space : $\mathcal{X} = [0, 1]^d$. Hence, here the dimensionality of the problem is d . The data will be uniformly sampled in the input space.
- output space : $\mathcal{Y} = \mathbb{R}$

— toy function to approximate : $f : x \in \mathbb{R}^d \mapsto \|x\|$ (euclidean norm).

Run a simulation that computes the test error of a kNN estimator, for different values of n and d , in order to observe the curse of dimensionality. The dataset should be uniformly sampled in the input space. You can experiment with the number of neighbors used of k , for instance 2.

You can use the template file `knn_uniform.py`. Note that as we know the function that we try to approximate, it is possible to perfectly evaluate the error made by the k-NN algorithm for each sample. For instance, you can observe images like the one in figure 3. In the simulation that generated the figure, the samples (both the dataset and the test samples) were drawn randomly in $\mathcal{X}$, with a uniform distribution (hence the dataset is not on a regular lattice, although you can do use a regular lattice in your simulation).




FIGURE 3 – Curse of dimensionality. We observe that the logarithm of the error is linear in the logarithm of the number of samples, with a slope that depends on the dimension d . When d is large, the problem is that the curve is very flat, which means that the error decreases very slowly.

1.4 Adaptivity

However, the situation is not always that catastrophic in high dimensions and we will illustrate this with an example. We keep exactly the same setting, but now the

input dataset X is given and fixed, and we use $d = 30$.

Perform a k -NN estimation using the data stored in `data_knn/` and use a varying number of samples n that you use from the dataset. Plot the error as a function of n .

You can use the template file `knn_adaptivity.py`. You should observe something like figure 4. We note that the learning rate is way better than $\mathcal{O}(n^{-1/d})$.

- What could be an explanation of this phenomenon?
- How could we test this hypothesis?




FIGURE 4 – Nearest neighbors estimation on a different dataset.