


# FTML practical session 6

3 juin 2026

Gradient descent: squared distance to the OLS estimator

$$\|\theta - \hat{\theta}\|^2$$



## TABLE DES MATIÈRES

|     |                                                              |   |
|-----|--------------------------------------------------------------|---|
| 1   | Gradient descent and line search for a least squares problem | 2 |
| 1.1 | Setting                                                      | 2 |
| 1.2 | Vanilla gradient descent                                     | 3 |
| 1.3 | Line search : a nested optimization problem                  | 3 |

## 1 GRADIENT DESCENT AND LINE SEARCH FOR A LEAST SQUARES PROBLEM

We come back to a least squares problem, that we will use to show an example of gradient descent algorithm, with a non-constant learning rate.

### 1.1 Setting

#### 1.1.1 Notations

We recall the setting and the notations here :

- $\mathcal{X} = \mathbb{R}^d$
- $\mathcal{Y} = \mathbb{R}$
- $\theta = \mathbb{R}^d$  defines the estimator, which is the function  $x \rightarrow x^T \theta$ .
- design matrix :  $X \in \mathbb{R}^{n,d}$
- vector of outputs (labels) :  $y \in \mathbb{R}^n$ .

We want to minimize the function  $f$  representing the empirical risk :

$$f(\theta) = \frac{1}{2n} \|X\theta - y\|^2 \quad (1)$$

The  $\frac{1}{2}$  is just convenient here because it disappears in the gradient.

#### 1.1.2 Gradient descent

We recall that the gradient of  $f$  and its Hessian write :

$$\begin{aligned} \nabla_{\theta} f(\theta) &= \frac{1}{n} X^T (X\theta - y) \\ &= H\theta - \frac{1}{n} X^T y \end{aligned} \quad (2)$$

$$H = \frac{1}{n} X^T X \quad (3)$$

An **iteration** of the gradient algorithm writes :

$$\theta_{t+1} = \theta_t - \gamma \nabla_{\theta} f(\theta_t) \quad (4)$$

where  $\gamma > 0$  is the learning rate.

#### 1.1.3 Global minimizers

We consider the global minimizers of  $f$ , noted  $\theta^*$ . They all necessary verify

$$\nabla_{\theta} f(\theta^*) = 0 \quad (5)$$

Which means that

$$H\theta^* = \frac{1}{n} X^T y \quad (6)$$

As  $\theta \rightarrow f(\theta)$  is convex, we know that this condition is also sufficient : any vector that cancels the gradient is a global minimum.

If  $H$  is not invertible, there might be several minimizers. However, if  $f$  is strongly convex, then  $H$  is invertible and  $\theta^*$  is unique, we note it  $\theta^*$ . We will consider such a case here. In general,  $H$  is a positive semi-definite matrix, and here we hence consider a case where it is definite positive.

### 1.2 Vanilla gradient descent

Implement a gradient descent (GD) on the data contained in `exercice_1_line_search/data/`, and experiment the learning rate in order to observe convergence or divergence.

Template files in the folder :

- `main.py`
- `algorithms.py` (to edit)

You should observe something like figure 1 in the cases of convergence.




FIGURE 1 – Constant step size gradient descent and for several learning rates.

### 1.3 Line search : a nested optimization problem

Considering a fixed iteration step  $\theta_t$ , we note

$$\alpha(\gamma) = \theta_t - \gamma \nabla_{\theta} f(\theta_t) \in \mathbb{R}^d \quad (7)$$


The **exact line search** method attempts to find the optimal step  $\gamma^*$ , at each iteration. This means, given the position  $\theta_t$ , the parameter  $\gamma$  that minimizes the function defined by

$$\begin{aligned} g(\gamma) &= f(\theta_t - \gamma \nabla_{\theta} f(\theta_t)) \\ &= f(\alpha(\gamma)) \end{aligned} \quad (8)$$

Is  $g : \mathbb{R} \rightarrow \mathbb{R}$  a convex function ?

Find the value  $\gamma^*$  that minimizes  $\gamma \rightarrow g(\gamma)$  for a given  $\theta_t$ , and perform gradient descent where  $\gamma$  is set optimally thanks to this method (exact line search). Compare the convergence speeds by measuring the distance between the iterate and the OLS estimator  $\hat{\theta}$  at each iteration, and experiment with the dimensions of the problem.

You can uncomment the lines that call the `line_search()` function in the main file. You should observe something like figure 2



**FIGURE 2** – Comparison between constant step size gradient descent and line search. In this image, we compare the number of iterations and running time to reach a value for  $\|\theta - \hat{\theta}\|$  where  $\hat{\theta}$  is the OLS estimator and  $\theta$  the iterate. The right plot is also in full-log scale, as opposed to the previous plots. We note that for these dimensions, OLS solved analytically is way faster than GD and line search. For way larger dimensions (typically  $d > 10^5$ ), the balance might be in favor of GD-type methods.




FIGURE 3 – If the learning rate is too large, we might have a divergence!




FIGURE 4 – Same image for larger dimensions of the problem.